



# Energy-Aware Computing in Heterogeneous Data Centers

## Terminology

Max Plauth, Sven Köhler, Felix Eberhardt, Lukas Wenzel and Andreas Polze  
Operating Systems and Middleware Group



### Chart 2

# Terminology

---

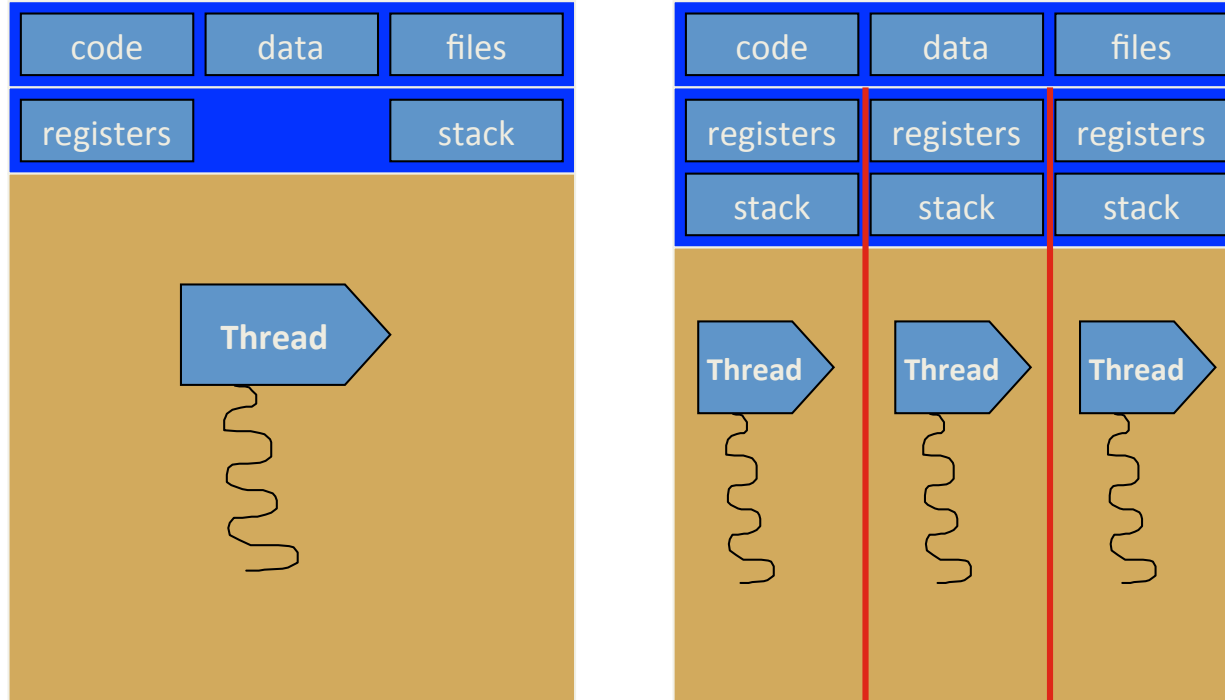
## ■ **Concurrency**

- Capability of a system to have two or more activities in progress at the same time
- May be independent, loosely coupled or closely coupled
- Classical operating system responsibility for a better utilization of CPU, memory, network, and other resources
- Demands scheduling and synchronization

## ■ **Parallelism**

- Capability of a system to execute activities simultaneously
  - Demands parallel hardware, concurrency support, (and communication)
- Any parallel program is a concurrent program
  - Some concurrent programs cannot be run as parallel program

# Example: Operating System



**Energy-Aware  
Computing  
Seminar**

Sven Köhler

Chart 4

# Concurrency Is Hard

---

- **Deadlock**

- Two or more processes / threads are unable to proceed
- Each is waiting for one of the others to do something

- **Livelock**

- Two or more processes / threads continuously change their states in response to changes in the other processes / threads
- No global progress for the application

- **Race condition**

- Two or more processes / threads are executed concurrently
- Final result of the application depends on the relative timing of their execution

# Race Condition

---

```
void echo() {  
    char_in = getchar();  
    char_out = char_in;  
    putchar(char_out);  
}
```

- One piece of code in one process, executed at the same time ...
  - ... by two threads on a single core.
  - ... by two threads on two cores.
- What happens ?

# Terminology

---

## ■ Starvation

- A runnable process / thread is overlooked indefinitely
- Although it is able to proceed, it is never chosen to run (dispatching / scheduling)


## ■ Atomic Operation

- Function or action implemented as a sequence of one or more instructions
- Appears to be indivisible - no other process / thread can see an intermediate state or interrupt the operation
- Executed as a group, or not executed at all

## ■ Mutual Exclusion

- The requirement that when one process / thread is using a resource, no other shall be allowed to do that

# The Parallel Programming Problem



**Energy-Aware  
Computing  
Seminar**

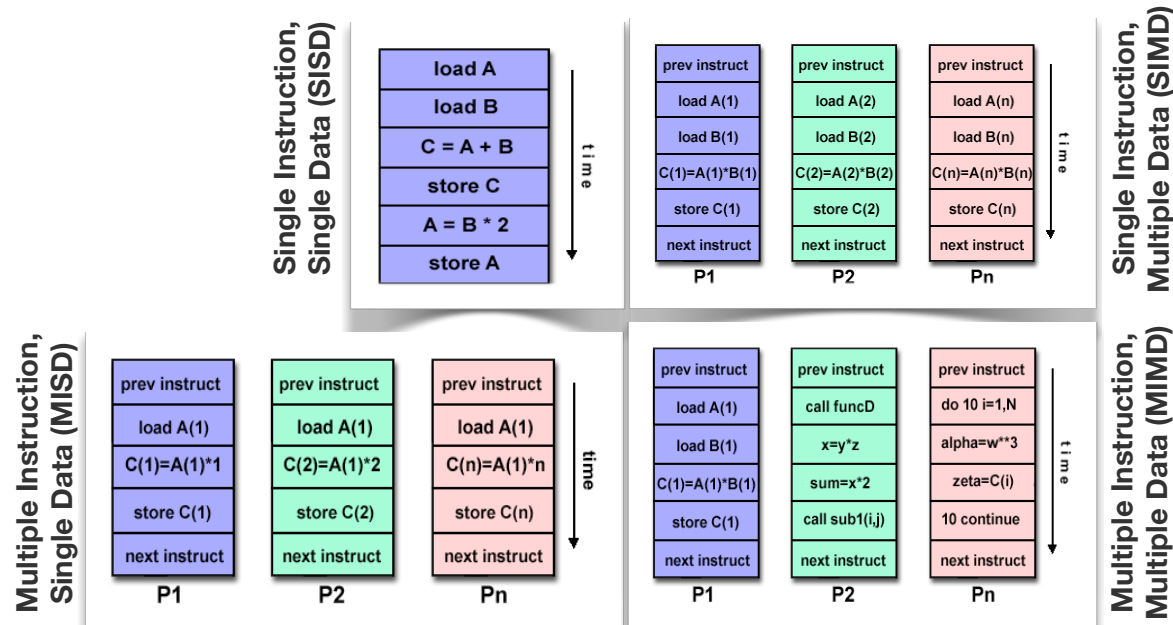
Sven Köhler

Chart 8





# Multiprocessor: Flynn's Taxonomy (1966)

- Classify multiprocessor architectures among **instruction** and **data** processing **dimension**



# Shared Memory vs. Shared Nothing


- Pfister: „shared memory“ vs. „distributed memory“
- Foster: „multiprocessor“ vs. „multicomputer“
- Tannenbaum: „shared memory“ vs. „private memory“



# Shared Memory vs. Shared Nothing

- Organization of parallel processing hardware as ...
  - **Shared memory system**
    - Concurrent processes can directly access a common address space
    - Typically implemented as memory hierarchy, with different cache levels
    - Examples: SMP systems, distributed shared memory systems, virtual runtime environment
  - **Shared nothing system**
    - Concurrent processes can only access local memory and exchange messages with other processes
    - Message exchange typically order of magnitudes slower than memory
    - Examples: Cluster systems, distributed systems (Hadoop, Grids, ...)

# The Parallel Programming Problem



**Energy-Aware  
Computing  
Seminar**

Sven Köhler

Chart **12**

# Programming Paradigm

---

- **Programming paradigm**
  - Coding convention or standard
  - Something a majority of people agrees upon
- **Parallel programming** is one of these paradigms
  - Other: Declarative, constraint-based, object-oriented
- Each paradigm can be realized by a set of **programming models**
- **Programming model:** „*set of rules for a game*“ [Almasi]
  - Programs and algorithms as game strategies
  - Point where execution environment and application meet
  - High-level view of the application on it's run time environment
  - Hardware might imply a model, but does not enforce it
  - For uniprocessor, no question due to „von Neumann“
  - Delivering performance while raising the level of abstraction